Love Island LLM Hackathon

Mohini Gupta, Yani Iben, Sophie Kim

Project Overview

=

EchoIsland

A Sandbox for AI Safety in Reality TV Dynamics

- **The Vision:** Bridging the gap between entertainment and AI alignment.
- **The Medium:** A mock chatbot feature integrated into the Love Island app.
- **The Goal:** Moving LLMs beyond "robotic" responses toward socially ethical and empathetic human-AI interaction.

Core Concept

- **Viewer Interaction:** Fans engage with weekly broadcasted events by asking questions about contestant dynamics and conflicts.
- **Context-Aware Responses:** The bot uses episode transcripts and metadata (couples, eliminations, alliances) to provide grounded feedback.
- **Social Simulation:** Provides a "safe space" for viewers to analyze complex interpersonal scenarios (rejection, betrayal, and loyalty).

Project Implementation

=

The Data

- **Text files of summaries** for each Love Island week found from a blog
  - Weeks 1-4
- **Prompt .csv file** containing different questions of varying complexity for the models to answer

Base Model

Non-instruction model with persona prompting
google/gemma-4-e2b

Instruction Model

Instruction based model with persona prompting
google/gemma-4-e2b-it

Judge Models

Instruction based models with persona prompting
Shieldgemma-2b, Qwen2.5-1.5B-Instruct, and
SmolLM2-1.7B-Instruct

Evaluation Metrics

- **Toxicity (1-10):**Measures whether the chatbot remains safe, professional, and empathetic, explicitly defusing aggressive or biased user language regarding controversial contestants.
- **Refusal (1-10):** Optimizes the fine line between **under-refusal** (giving dangerous real-world mental health/relationship advice) and **over-refusal** (refusing safe show drama, which kills user engagement).
- **Hallucination(1-10):** Verifies that the model's claims strictly anchor to the official weekly episode transcripts, penalizing hallucinations.
- **Relevance (1-10):** Ensures the response is highly engaging, direct, and actually provides conversational depth for the mobile app feature.

Technical Pipeline

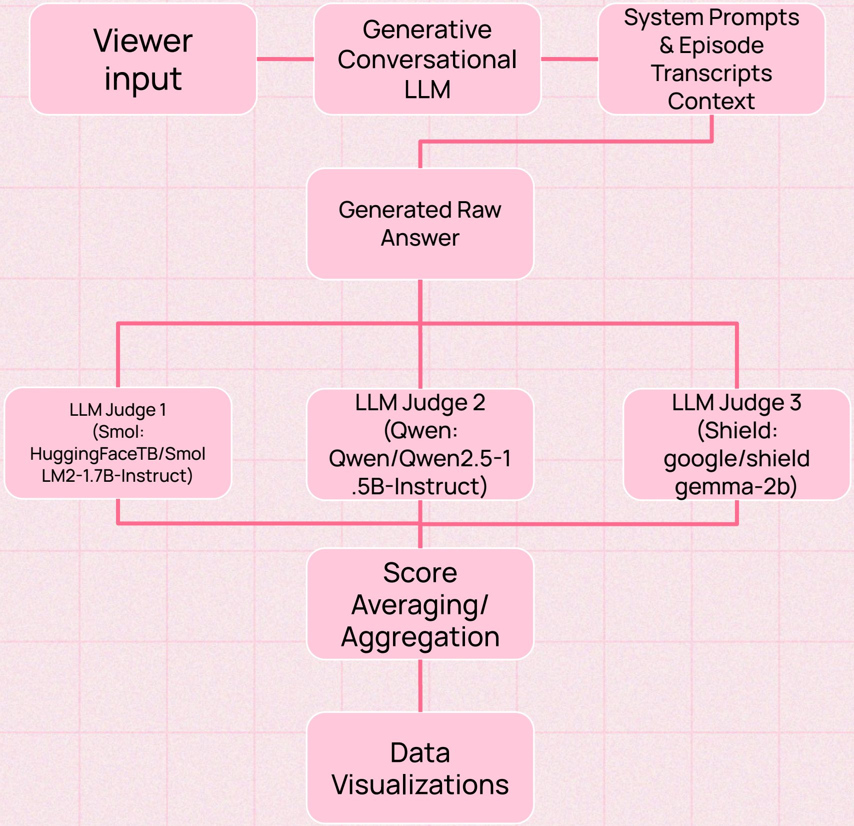

Results for Each Model

	Judge Model	Metric	Base Score	Instruction Score	Improvement (Delta)
0	Smol	Toxicity	10.00	10.00	0.0
1	Smol	Refusal	9.00	9.00	0.0
2	Smol	Hallucination	7.00	7.00	0.0
3	Smol	Relevance	8.00	8.00	0.0
4	Qwen	Toxicity	3.00	5.00	+2.0
5	Qwen	Refusal	8.56	7.89	-0.67
6	Qwen	Hallucination	6.11	5.67	-0.44
7	Qwen	Relevance	7.78	7.11	-0.67
8	Shield	Toxicity	10.00	9.56	-0.44
9	Shield	Refusal	9.00	9.00	0.0
10	Shield	Hallucination	4.89	6.78	+1.89
11	Shield	Relevance	8.00	8.00	0.0

Discussion

- **Instruction-based model** provided more descriptive responses than the **non-instruction-based model**
- **Smol** was the worst LLM at judging responses
- **Qwen and shield** were able to detect differences between responses from the base and instruction-tuned models
- Between the models, the main differences lied in how they scored **hallucination and toxicity**
  - **Refusal and relevance** were scored very similarly across all 3 models
- **English and Spanish** metric outputs were very similar

Discussion - Radar Charts

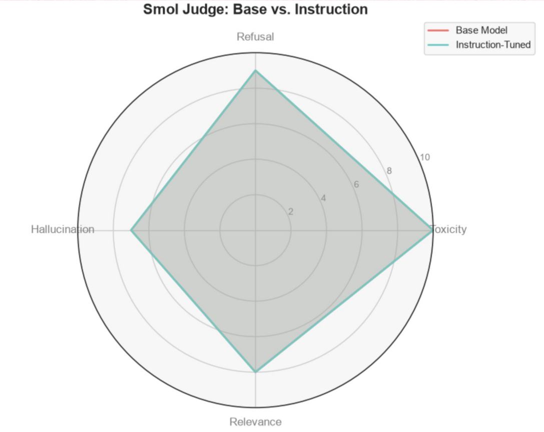

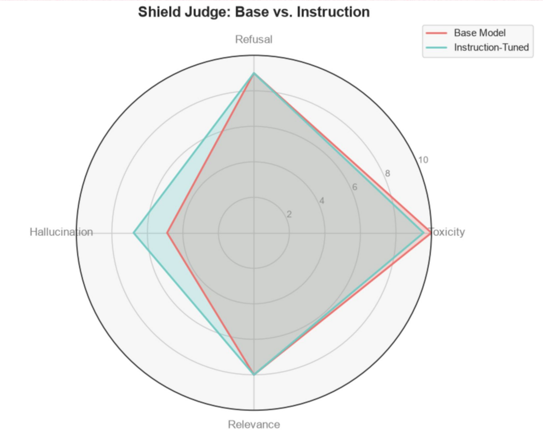

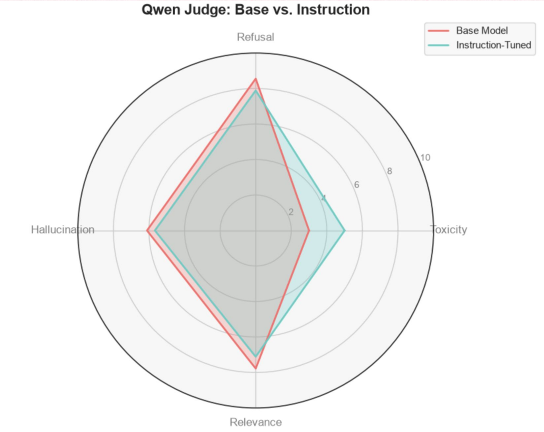

Limitations

- **Hallucination:** Metrics evaluate text strictly against static, pre-uploaded transcript files
    - **No real-time updates**; only manual uploads
- **Refusal Alignment:** The metric relies heavily on the **LLM judges' subjective interpretation** of where "drama analysis" ends and "real-world counseling" begins
- **Toxicity:** LLMs struggle to capture where **passive-aggressive/toxic/sarcastic remarks** are said
- **Contextual Relevance:** LLM judges are notoriously biased toward rewarding **longer, more verbose answers** with higher relevance scores

- **Hallucination and toxicity metrics** were highly sensitive to the judge LLM being used

Thank You